

Question: What step should we take to improve and make the model more efficient ??

1. Improve Context Quality

Instead of sending:

Top 5 chunks

Send:

Section Name

Page Number

Court Name

Chunk Content

Structured context helps the model reason better.

Example:

```
{  
  "court": "Supreme Court",  
  "section": "Final Judgment",  
  "page": 12,  
  "content": "..."  
}
```

2. Better Prompt Engineering

Avoid generic prompts.

Bad:

Answer the question.

Better:

You are a legal assistant.

Only answer using the provided context.

If the answer is not present in the context,
respond with:

"Insufficient information available."

Do not make assumptions.

This alone can reduce hallucinations significantly.

3. Add Source Citations

Make the model cite where information came from.

Example:

According to Paragraph 18 of the judgment...

or

Source: Page 12, Paragraph 4

Users trust answers more and you can verify them.

Currently, NyaySaar is still in its early stage, so our main focus is improving the accuracy of the responses. Before jumping into fine tuning, we want to strengthen the RAG pipeline by improving retrieval quality, chunking, and prompt engineering. We believe that better retrieval and better prompts will help reduce hallucinations significantly, and once we reach a good baseline, we can fine tune the model on legal data to make it more reliable for real world use.

On the product side, we are planning to improve the dashboard and introduce different user personas such as the general public and legal advisors, so that the experience can be tailored according to the user's needs. Right now, the platform mainly summarizes court orders, but in the next phase we want it to provide more meaningful legal assistance by analyzing the current status of a case and helping users understand possible next steps.

At the moment, the system only works with PDF court orders. Going forward, we plan to add OCR support so that scanned documents can also be processed. We also want to support regional languages like Hindi and Gujarati to make the platform more accessible. From a scalability perspective, we are looking at adding Redis for caching and implementing proper monitoring and logging so that the system remains fast, reliable, and easier to maintain as the number of users grows.